ANTHROP\C English

Explore features › Batch processing

EXPLORE FEATURES

Batch processing

Batch processing is a powerful approach for handling large volumes of requests efficiently. Instead of processing requests one at a time with immediate responses, batch processing allows you to submit multiple requests together for asynchronous processing. This pattern is particularly useful when:

- You need to process large volumes of data
- Immediate responses are not required
- You want to optimize for cost efficiency
- You're running large-scale evaluations or analyses

The Message Batches API is our first implementation of this pattern.

Message Batches API

The Message Batches API is a powerful, cost-effective way to asynchronously process large volumes of **Messages** requests. This approach is well-suited to tasks that do not require immediate responses, with most batches finishing in less than 1 hour while reducing costs by 50% and increasing throughput.

You can **explore the API reference directly**, in addition to this guide.

How the Message Batches API works

3. You can poll for the status of the batch and retrieve results when processing has ended for all requests.

This is especially useful for bulk operations that don't require immediate results, such as:

Large-scale evaluations: Process thousands of test cases efficiently.

Content moderation: Analyze large volumes of user-generated content asynchronously.

Data analysis: Generate insights or summaries for large datasets.

Bulk content generation: Create large amounts of text for various purposes (e.g., product descriptions, article summaries).

Batch limitations

A Message Batch is limited to either 100,000 Message requests or 256 MB in size, whichever is reached first.

We process each batch as fast as possible, with most batches completing within 1 hour. You will be able to access batch results when all messages have completed or after 24 hours, whichever comes first. Batches will expire if processing does not complete within 24 hours.

Batch results are available for 29 days after creation. After that, you may still view the Batch, but its results will no longer be available for download.

Claude 3.7 Sonnet supports up to 128K output tokens using the **extended output capabilities**.

Batches are scoped to a **Workspace**. You may view all batches—and their results—that were created within the Workspace that your API key belongs to.

Rate limits apply to both Batches API HTTP requests and the number of requests within a batch waiting to be processed. See **Message Batches API rate limits**. Additionally, we

ANTHROP\C

Supported models

The Message Batches API currently supports:

- Claude 3.7 Sonnet (`claude-3-7-sonnet-20250219`)
- Claude 3.5 Sonnet (`claude-3-5-sonnet-20240620` and `claude-3-5-sonnet-20241022`)
- Claude 3.5 Haiku (`claude-3-5-haiku-20241022`)
- Claude 3 Haiku (`claude-3-haiku-20240307`)
- Claude 3 Opus (`claude-3-opus-20240229`)

What can be batched

Any request that you can make to the Messages API can be included in a batch. This includes:

- Vision
- Tool use
- System messages
- Multi-turn conversations
- Any beta features

Since each request in the batch is processed independently, you can mix different types of requests within a single batch.

Pricing

ANTHROP\C

Claude 3.7 Sonnet	$1.50 / MTok	$7.50 / MTok
Claude 3.5 Sonnet	$1.50 / MTok	$7.50 / MTok
Claude 3.5 Haiku	$0.40 / MTok	$2 / MTok
Claude 3 Opus	$7.50 / MTok	$37.50 / MTok
Claude 3 Haiku	$0.125 / MTok	$0.625 / MTok

How to use the Message Batches API

Prepare and create your batch

A Message Batch is composed of a list of requests to create a Message. The shape of an individual request is comprised of:

- A unique `custom_id` for identifying the Messages request
- A `params` object with the standard Messages API parameters

You can create a batch by passing this list into the `requests` parameter:

Shell | Python | TypeScript | Java

```
curl https://api.anthropic.com/v1/messages/batches \
     --header "x-api-key: $ANTHROPIC_API_KEY" \
     --header "anthropic-version: 2023-06-01" \
     --header "content-type: application/json" \
     --data \
'{
    "requests": [
        {
            "custom_id": "my-first-request",
```

ANTHROP\C

Explore features › **Batch processing**

```
                ]
            }
        },
        {
            "custom_id": "my-second-request",
            "params": {
                "model": "claude-3-7-sonnet-20250219",
                "max_tokens": 1024,
                "messages": [
                    {"role": "user", "content": "Hi again, friend"}
                ]
            }
        }
    ]
}'
```

In this example, two separate requests are batched together for asynchronous processing. Each request has a unique `custom_id` and contains the standard parameters you'd use for a Messages API call.

Test your batch requests with the Messages API

Validation of the `params` object for each message request is performed asynchronously, and validation errors are returned when processing of the entire batch has ended. You can ensure that you are building your input correctly by verifying your request shape with the Messages API first.

When a batch is first created, the response will have a processing status of `in_progress`.

JSON

ANTHROP\C


```
  "request_counts": {
    "processing": 2,
    "succeeded": 0,
    "errored": 0,
    "canceled": 0,
    "expired": 0
  },
  "ended_at": null,
  "created_at": "2024-09-24T18:37:24.100435Z",
  "expires_at": "2024-09-25T18:37:24.100435Z",
  "cancel_initiated_at": null,
  "results_url": null
}
```


Tracking your batch

The Message Batch's `processing_status` field indicates the stage of processing the batch is in. It starts as `in_progress`, then updates to `ended` once all the requests in the batch have finished processing, and results are ready. You can monitor the state of your batch by visiting the Console, or using the retrieval endpoint:

Shell | Python | TypeScript | Java

```
curl https://api.anthropic.com/v1/messages/batches/msgbatch_01HkcTjaV5uDC8jWR4ZsD
 --header "x-api-key: $ANTHROPIC_API_KEY" \
 --header "anthropic-version: 2023-06-01" \
 | sed -E 's/.*"id":"([^"]+)".*"processing_status":"([^"]+)".*/Batch \1 processin
```

You can poll this endpoint to know when processing has ended.

Retrieving batch results

ANTHROP\C

`succeeded`	Request was successful. Includes the message result.
`errored`	Request encountered an error and a message was not created. Possible errors include invalid requests and internal server errors. You will not be billed for these requests.
`canceled`	User canceled the batch before this request could be sent to the model. You will not be billed for these requests.
`expired`	Batch reached its 24 hour expiration before this request could be sent to the model. You will not be billed for these requests.

You will see an overview of your results with the batch's `request_counts`, which shows how many requests reached each of these four states.

Results of the batch are available for download at the `results_url` property on the Message Batch, and if the organization permission allows, in the Console. Because of the potentially large size of the results, it's recommended to **stream results** back rather than download them all at once.

Shell | Python | TypeScript | Java

```shell
#!/bin/sh
curl "https://api.anthropic.com/v1/messages/batches/msgbatch_01HkcTjaV5uDC8jWR4Zs
  --header "anthropic-version: 2023-06-01" \
  --header "x-api-key: $ANTHROPIC_API_KEY" \
  | grep -o '"results_url":[[:space:]]*"[^"]*"' \
  | cut -d'"' -f4 \
  | while read -r url; do
    curl -s "$url" \
      --header "anthropic-version: 2023-06-01" \
      --header "x-api-key: $ANTHROPIC_API_KEY" \
      | sed 's/}{/}\n{/g' \
      | while IFS= read -r line
    do
      result_type=$(echo "$line" | sed -n 's/.*"result":[[:space:]]*{[[:space:]]*
      custom_id=$(echo "$line" | sed -n 's/.*"custom_id":[[:space:]]*"\([^"]*\)".
```

ANTHROP\C

Explore features › **Batch processing**

```
          ;;
        "errored")
          if [ "$error_type" = "invalid_request" ]; then
            # Request body must be fixed before re-sending request
            echo "Validation error: $custom_id"
          else
            # Request can be retried directly
            echo "Server error: $custom_id"
          fi
          ;;
        "expired")
          echo "Expired: $line"
          ;;
      esac
    done
  done
```

The results will be in `.jsonl` format, where each line is a valid JSON object representing the result of a single request in the Message Batch. For each streamed result, you can do something different depending on its `custom_id` and result type. Here is an example set of results:

.jsonl file


```
{"custom_id":"my-second-request","result":{"type":"succeeded","message":{"id":"ms
{"custom_id":"my-first-request","result":{"type":"succeeded","message":{"id":"msg
```


If your result has an error, its `result.error` will be set to our standard error shape.

Batch results may not match input order

ANTHROP\C

Using prompt caching with Message Batches

The Message Batches API supports prompt caching, allowing you to potentially reduce costs and processing time for batch requests. The pricing discounts from prompt caching and Message Batches can stack, providing even greater cost savings when both features are used together. However, since batch requests are processed asynchronously and concurrently, cache hits are provided on a best-effort basis. Users typically experience cache hit rates ranging from 30% to 98%, depending on their traffic patterns.

To maximize the likelihood of cache hits in your batch requests:

1. Include identical `cache_control` blocks in every Message request within your batch
2. Maintain a steady stream of requests to prevent cache entries from expiring after their 5-minute lifetime
3. Structure your requests to share as much cached content as possible

Example of implementing prompt caching in a batch:

Shell | Python | TypeScript | Java

```
curl https://api.anthropic.com/v1/messages/batches \
     --header "x-api-key: $ANTHROPIC_API_KEY" \
     --header "anthropic-version: 2023-06-01" \
     --header "content-type: application/json" \
     --data \
'{
    "requests": [
        {
            "custom_id": "my-first-request",
            "params": {
                "model": "claude-3-7-sonnet-20250219",
```

ANTHROP\C

Explore features › **Batch processing**

```
                    },
                    {
                        "type": "text",
                        "text": "<the entire contents of Pride and Prejudice>",
                        "cache_control": {"type": "ephemeral"}
                    }
                ],
                "messages": [
                    {"role": "user", "content": "Analyze the major themes in Prid
                ]
            }
        },
        {
            "custom_id": "my-second-request",
            "params": {
                "model": "claude-3-7-sonnet-20250219",
                "max_tokens": 1024,
                "system": [
                    {
                        "type": "text",
                        "text": "You are an AI assistant tasked with analyzing li
                    },
                    {
                        "type": "text",
                        "text": "<the entire contents of Pride and Prejudice>",
                        "cache_control": {"type": "ephemeral"}
                    }
                ],
                "messages": [
                    {"role": "user", "content": "Write a summary of Pride and Pre
                ]
            }
        }
    ]
}'
```

ANTHROP\C

Best practices for effective batching

To get the most out of the Batches API:

- Monitor batch processing status regularly and implement appropriate retry logic for failed requests.
- Use meaningful `custom_id` values to easily match results with requests, since order is not guaranteed.
- Consider breaking very large datasets into multiple batches for better manageability.
- Dry run a single request shape with the Messages API to avoid validation errors.

Troubleshooting common issues

If experiencing unexpected behavior:

- Verify that the total batch request size doesn't exceed 256 MB. If the request size is too large, you may get a 413 `request_too_large` error.
- Check that you're using supported models for all requests in the batch.
- Ensure each request in the batch has a unique `custom_id`.
- Ensure that it has been less than 29 days since batch `created_at` (not processing `ended_at`) time. If over 29 days have passed, results will no longer be viewable.
- Confirm that the batch has not been canceled.

Note that the failure of one request in a batch does not affect the processing of other requests.

ANTHROP\C

permission to view Workspace batches in the Console.

Result availability: Batch results are available for 29 days after the batch is created, allowing ample time for retrieval and processing.

FAQ

- How long does it take for a batch to process?
- Is the Batches API available for all models?
- Can I use the Message Batches API with other API features?
- How does the Message Batches API affect pricing?
- Can I update a batch after it's been submitted?
- Are there Message Batches API rate limits and do they interact with the Messages API rate limits?
- How do I handle errors in my batch requests?
- How does the Message Batches API handle privacy and data separation?
- Can I use prompt caching in the Message Batches API?
- How do I use beta features in the Message Batches API?

ANTHROP\C

Was this page helpful? Yes No

‹ Prompt caching

Extended thinking ›